



Exp A1: Request Latency Distribution by Block Length

Larger blocks reduce per-request steps but increase variance


Block Length = 16 tokens | CV = 0.270



Block Length = 32 tokens | CV = 0.270



Block Length = 64 tokens | CV = 0.270



Block Length = 128 tokens | CV = 0.271

